

Image Manipulation & Compression with Linear Algebra

tinyimg.amirghofran.com

Image Transforms

Rotation -28°



☐ Link Scale X/Y

Scale X 0.45x



Scale Y 0.50x



Shear X -0.49



Shear Y -0.40



Translate X 0 px




Translate Y 0 px



Flip Horizontal (Off)

Flip Vertical (Off)



Processing & Export

Filters

Blur

Sharpen

Edge

Emboss

SVD Compression

Rank 50



Apply SVD

Transformed Area

0.7236

Formula: $\text{Area} = 4 \times |\det(M)|$

Transformation Matrices

Overall Translate Rotate

Scale Shear

0.482	-0.451	0.000	0.000
0.052	0.326	0.000	0.000
0.000	0.000	1.000	0.000
0.000	0.000	0.000	1.000

Download Image

View on Github

Overview

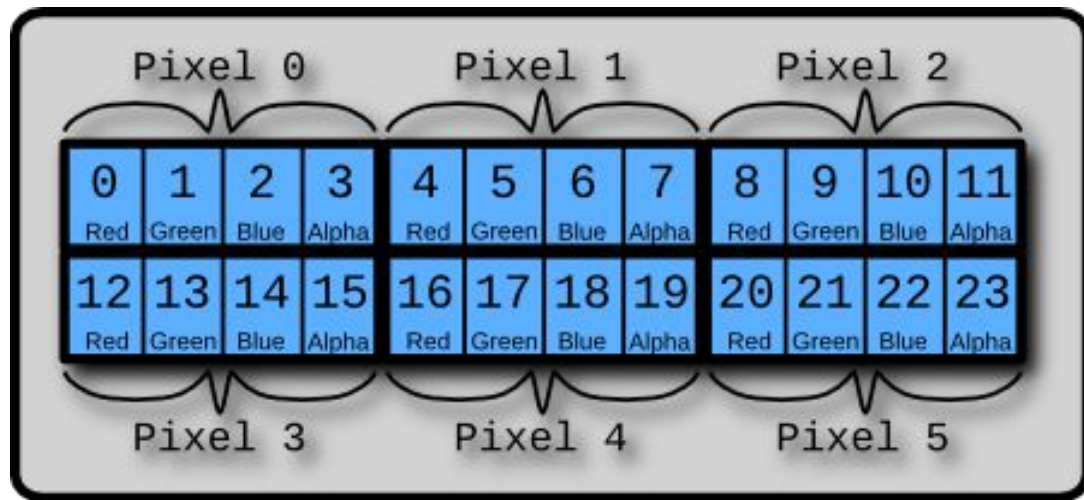
- Basic app for image manipulation
- Transformations: Rotation, Scale, Shear, Translate
- Filters: Blur, Sharpen, Edge Detection, Emboss
- Image Compression
- WebGL for image transformations
- Go logic for applying filters and image compression
- Compiled logic to WASM and loaded in the browser

Matrix Representation of Images

When an image is uploaded each pixel is represented with 4 values following the RGBA color model

Instead of storing each color channel as four separate matrices we start with a Clamped Array, a flat single-dimensional array. This ensures:

- Memory efficiency
- Fast data transfer & compatibility




Matrix Representation of Images: Processing

The backend splits the image into **4 color channels** — Red, Green, Blue, and Alpha — and turns each one into a **2D matrix**.

For an image that's **W×H** each channel becomes a matrix of size **H×W**, where each value represents the **intensity** of that channel at a specific pixel.

$$M_{\text{channel}} = \begin{pmatrix} p_{0,0} & p_{1,0} & \cdots & p_{W-1,0} \\ p_{0,1} & p_{1,1} & \cdots & p_{W-1,1} \\ \vdots & \vdots & \ddots & \vdots \\ p_{0,H-1} & p_{1,H-1} & \cdots & p_{W-1,H-1} \end{pmatrix}$$



Once we have these matrices, we can apply **linear algebra techniques** like:

- **Convolution**
- **SVD**

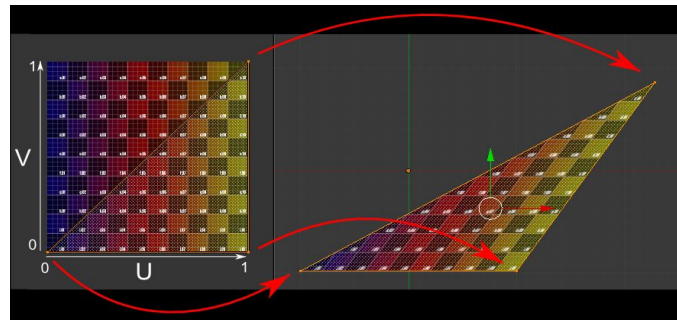
Arrays stacked over each other
to form a Digital Image.

Matrix Representation of Images: Transformations

The image's **pixel data** (RGBA values) are **not changed**, only **how the image appears on screen** — its position, scale, rotation, and shape. To do so we follow the following steps:

1. A **quadrilateral** is drawn to fill the screen made up of **4 vertices**. These define **where** the image should be drawn.
2. WebGL requires the vertices to be in 4D (**homogeneous coordinates**). So we convert each 2D vertex into a 4D vector
3. These 2 extra coordinates are called **UV (texture) coordinates**, which map the image to its location in the **quad**.
4. Now the transformations can be applied to the two position coordinates, transforming the **quad** only, and the image is painted onto this.

$$\vec{v} = \begin{bmatrix} x \\ y \\ 0 \\ 1 \end{bmatrix}$$



Linear Transformations

Linear Transformation

A function that transforms vectors while keeping their structure.

- $T(\mathbf{u} + \mathbf{v}) = T(\mathbf{u}) + T(\mathbf{v})$

- $T(k\mathbf{u}) = kT(\mathbf{u})$

Matrix Form of Transformation

All linear transformations can be written as $T(\mathbf{x}) = A\mathbf{x}$

$$A = \begin{bmatrix} 2 & -1 \\ 1 & 3 \end{bmatrix}, \quad \mathbf{x} = \begin{bmatrix} x \\ y \end{bmatrix} \quad \Rightarrow \quad T(\mathbf{x}) = \begin{bmatrix} 2x - y \\ x + 3y \end{bmatrix}$$

Rotation



Rotation Matrix:

$$\begin{bmatrix} \cos \theta & -\sin \theta \\ \sin \theta & \cos \theta \end{bmatrix}$$

- Counterclockwise rotation by default
- Rotates every pixel of the image around the center

Equation for rotation:

$$\text{rotate}(30^\circ) = \begin{bmatrix} \cos(30^\circ) & -\sin(30^\circ) \\ \sin(30^\circ) & \cos(30^\circ) \end{bmatrix} \begin{bmatrix} 0.866 & -0.5 \\ 0.5 & 0.866 \end{bmatrix}$$



Scaling


Scaling Matrix:

$$\begin{bmatrix} s_x & 0 \\ 0 & s_y \end{bmatrix}$$

- Scales horizontally by s_x and vertically by s_y
- Zoom in/out

Equation for scaling:

$$\text{scale}(2, 0.5) = \begin{bmatrix} 2 & 0 \\ 0 & 0.5 \end{bmatrix}$$



Shearing

Shearing Matrix:

Horizontal: $\begin{bmatrix} 1 & k \\ 0 & 1 \end{bmatrix}$



Vertical: $\begin{bmatrix} 1 & 0 \\ k & 1 \end{bmatrix}$



- Distorts the shape diagonally
- K controls how slanted the shear is

Equation for shearing:

$$\text{shear}_x(1.5) = \begin{bmatrix} 1 & 1.5 \\ 0 & 1 \end{bmatrix}$$

Combining Transformations = Matrix Multiplication

$$T = R \cdot S = \begin{bmatrix} 0 & -1 \\ 1 & 0 \end{bmatrix} \cdot \begin{bmatrix} 2 & 0 \\ 0 & 0.5 \end{bmatrix} \begin{bmatrix} 0 & -0.5 \\ 2 & 0 \end{bmatrix}$$

Calculating Area with Determinant

Determinant

- In the project, we apply transformations like rotation, scaling, and shearing using 4×4 matrices.
- These transformations affect how the image is displayed on screen.
- The determinant of the matrix plays a key role in understanding how the image's area changes.

Determinants

- Determinants help us understand how a transformation affects area.

For a 2D matrix:

$$\det \begin{pmatrix} a & c \\ b & d \end{pmatrix} = ad - bc$$

- $\det = ad - bc$
- The absolute value of the determinant tells us the area scaling factor.
- Possible values:
 - $|\det| = 1$: Area unchanged (ex. rotation)
 - $|\det| > 1$: Area increases (ex. scaling up)
 - $0 < |\det| < 1$: Area decreases (ex. slight compression)
 - $|\det| = 0$: Area collapses to a line or point (degenerate transformation)

How it works:

In TinyIMG, we use 4x4 matrices, but only the top-left 2x2 submatrix affects area:

$$\det \begin{pmatrix} a & c & 0 & tx \\ b & d & 0 & ty \\ 0 & 0 & 1 & 0 \\ 0 & 0 & 0 & 1 \end{pmatrix} = ad - bc$$


Area scaling factor = $|ad - bc|$

This value tells us how much larger or smaller the transformed image appears.

Filters with Kernel Convolutions

Convolution Basics


Convolution is a fundamental technique where a filter (or kernel) is applied to an image to modify its pixel values based on a weighted average of their neighbours.



Convolution Calculation

The convolution operation for a pixel at (x, y) and a channel c is mathematically defined as:

$$P'_{\text{out}}(x, y, c) = \sum_{i=-k}^k \sum_{j=-k}^k K(i, j) \cdot P_{\text{in}}(x + i, y + j, c)$$



Blur

Averages neighbor pixels.

$$K_{\text{blur}} = \frac{1}{9} \begin{pmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{pmatrix}$$



Sharpen

Enhances edges by emphasizing differences with neighbors.


$$K_{\text{sharpen}} = \begin{pmatrix} 0 & -1 & 0 \\ -1 & 5 & -1 \\ 0 & -1 & 0 \end{pmatrix}$$



Edge Detection

Highlights areas of rapid intensity change.

$$K_{\text{edge}} = \begin{pmatrix} -1 & -1 & -1 \\ -1 & 8 & -1 \\ -1 & -1 & -1 \end{pmatrix}$$



Emboss

Creates a raised or lowered effect based on edge direction.

$$K_{\text{emboss}} = \begin{pmatrix} -2 & -1 & 0 \\ -1 & 1 & 1 \\ 0 & 1 & 2 \end{pmatrix}$$






Compression with SVD



Separating Different Color Channels



Pixel of an RGB image are formed from the corresponding pixel of the three component images

Calculating Singular Value Decomposition $A = U \Sigma V^T$

- **Left Singular Vectors (U):** The columns of the matrix U are the eigenvectors of the matrix AA^T . These vectors form an orthonormal basis for the codomain space.
- **Singular Values (Σ):** The diagonal entries of the matrix S are the singular values of A , arranged in descending order. As mentioned above, these singular values (σ_i) are the square roots of the eigenvalues (λ_i) of $A^T A$ (or AA^T).

$$\sigma_i = \sqrt{\lambda_i(A^T A)} = \sqrt{\lambda_i(AA^T)}$$

- **Right Singular Vectors (V):** The columns of the matrix V are the eigenvectors of the matrix $A^T A$. These vectors form an orthonormal basis for the domain space.

Calculating Singular Value Decomposition

$$A = U \Sigma V^T$$

$$\mathbf{A}_{m \times n} = \left[\begin{array}{c|c|c|c} \vec{u}_1 & \vec{u}_2 & \dots & \vec{u}_m \end{array} \right] \begin{bmatrix} \sigma_1 & 0 & \dots & 0 \\ 0 & \sigma_2 & \dots & 0 \\ \vdots & \vdots & \ddots & \vdots \\ 0 & 0 & \dots & \sigma_{\min(m,n)} \\ \vdots & \vdots & & \vdots \\ 0 & 0 & \dots & 0 \end{bmatrix}$$

$m \times m$

$m \times n$

$$\left[\begin{array}{c|c|c|c} \vec{v}_1 & \vec{v}_2 & \dots & \vec{v}_n \end{array} \right]^T$$

$n \times n$

</>

</>

Reconstruction with k ranks

We reconstruct the approximation image (matrix) by multiplying the SVD decomposition components.

$$A_k = U_k \Sigma_k V_k^T.$$

$$\begin{array}{ccc} \mathbf{A}_k \approx \left[\begin{array}{c|c|c|c} \vec{u}_1 & \vec{u}_2 & \dots & \vec{u}_k \end{array} \right] & \begin{bmatrix} \sigma_1 & 0 & \dots & 0 \\ 0 & \sigma_2 & \dots & 0 \\ \vdots & \vdots & \ddots & \vdots \\ 0 & 0 & \dots & \sigma_k \end{bmatrix} & \left[\begin{array}{c|c|c|c} \vec{v}_1 & \vec{v}_2 & \dots & \vec{v}_k \end{array} \right]^T \\ m \times k & k \times k & k \times n \end{array}$$

Original



Rank 70



Rank 40



Rank 10



Challenges

Compile code to
WebAssembly

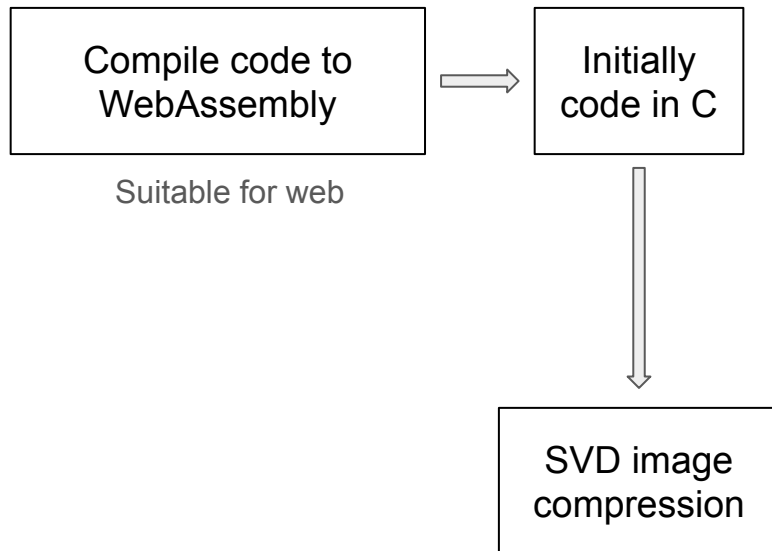
Suitable for web

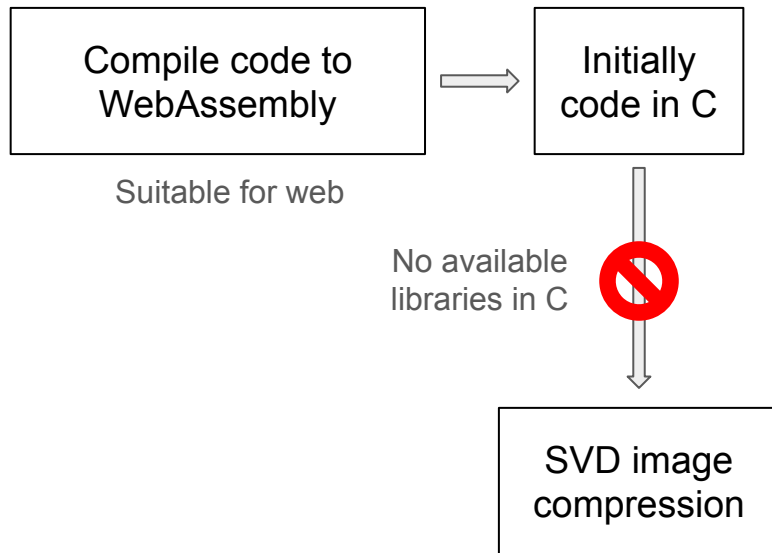
Compile code to
WebAssembly

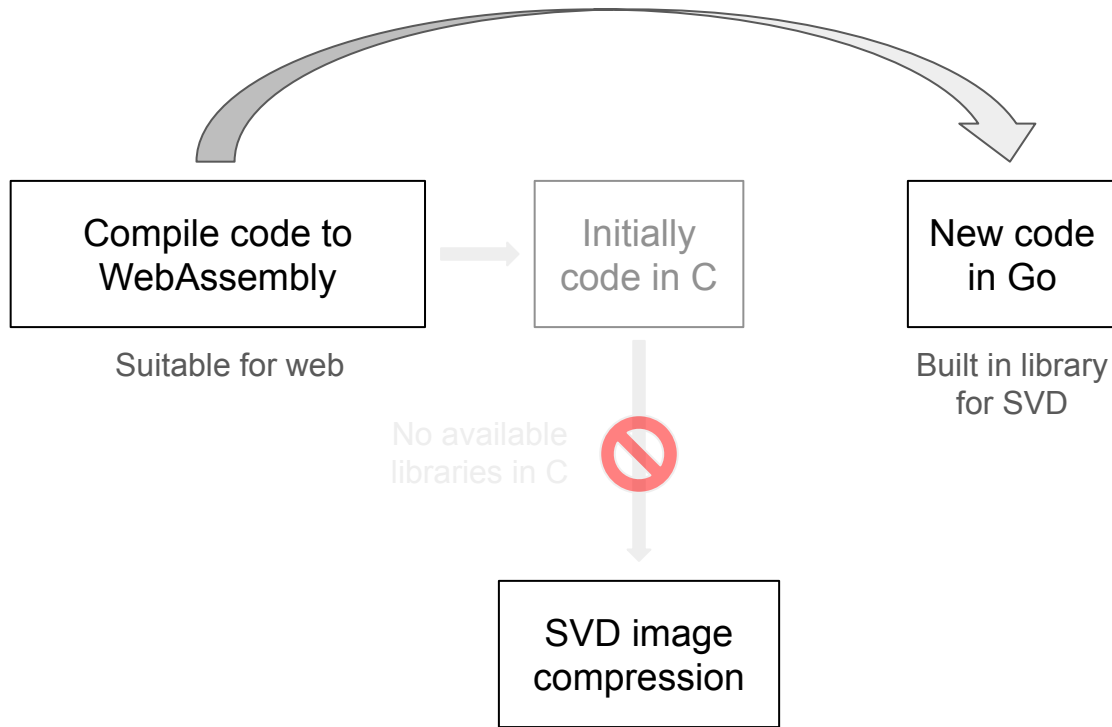


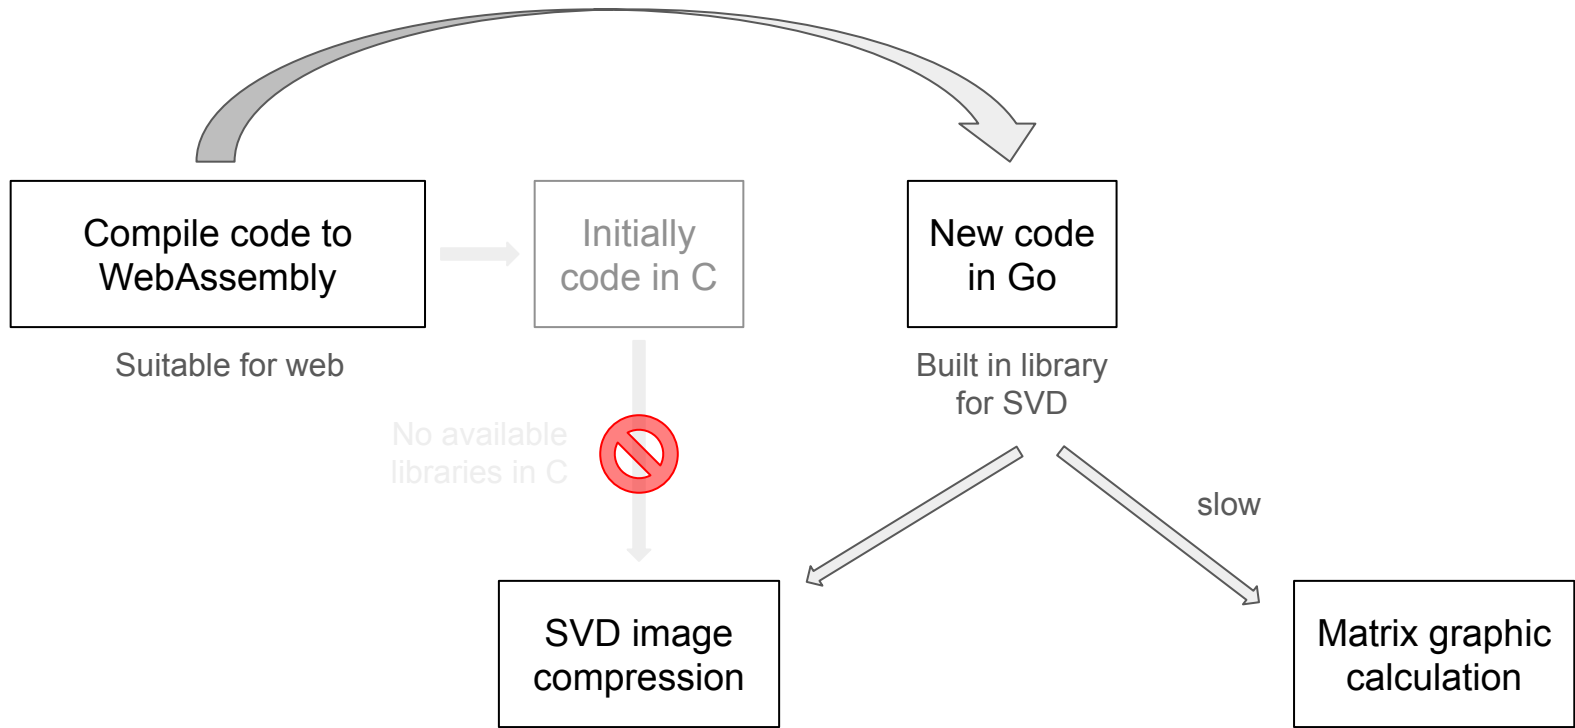
Initially
code in C

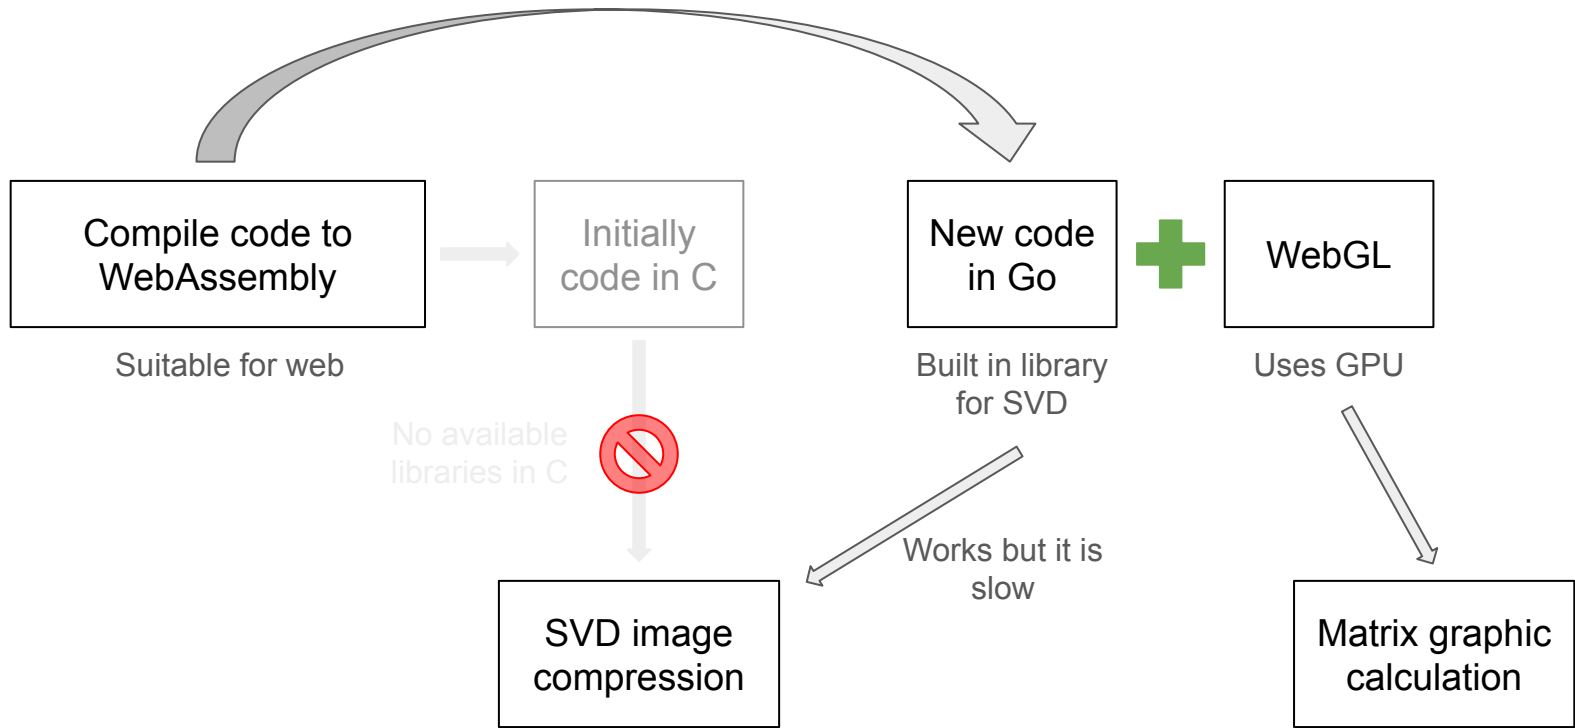
Suitable for web











Thank you